

Assignment 1 - DSCI-6601-001 (Pract Machine Learning 77223)

January 15, 2023

0.1 Name: Daniel Kojo Afealete Fiadjoe

0.2 Student ID: 202291439

0.3 Course: DSCI-6601-001 (Pract Machine Learning 77223)

1 Assignment 1

1.0.1 Import Libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

# Libraries to split data, create simple linear regression
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.metrics import mean_squared_error
```

1.1 Question 2

Write some Python code that can load in a .CSV file. Load in the Heart.csv file from the “Introduction to Statistical Learning” website (<https://www.statlearning.com/s/Heart.csv>). Load the data into a numpy array and have the code print the column headings in the dataset and the 25th data row/observation contained in the file.

```
[2]: # Code to load in the Heart.csv and show the rows and columns
data = pd.read_csv('https://www.statlearning.com/s/Heart.csv')

# Code to Load the dataset into numpy array called datanp
datanp = np.array(data)
```

```
[3]: # Code to display the number of rows and columns
print(f"There are {data.shape[0]} rows and {data.shape[1]} columns.")
```

There are 303 rows and 15 columns.

```
[4]: # Code to display first 5 rows of the data.
data.head()
```

```
[4]: Unnamed: 0  Age  Sex      ChestPain  RestBP  Chol  Fbs  RestECG  MaxHR  \
0           1   63    1      typical      145   233   1         2    150
1           2   67    1  asymptomatic      160   286   0         2    108
2           3   67    1  asymptomatic      120   229   0         2    129
3           4   37    1   nonanginal      130   250   0         0    187
4           5   41    0   nontypical      130   204   0         2    172

      ExAng  Oldpeak  Slope  Ca      Thal  AHD
0         0       2.3     3  0.0    fixed  No
1         1       1.5     2  3.0    normal  Yes
2         1       2.6     2  2.0  reversable  Yes
3         0       3.5     3  0.0    normal  No
4         0       1.4     1  0.0    normal  No
```

```
[5]: # Code to print all the column headings in the dataset.
for col in data.columns:
    print (col)
```

```
Unnamed: 0
Age
Sex
ChestPain
RestBP
Chol
Fbs
RestECG
MaxHR
ExAng
Oldpeak
Slope
Ca
Thal
AHD
```

```
[6]: # Print the 25th data row in the dataset
data.loc[25]
```

```
[6]: Unnamed: 0      26
Age      50
Sex      0
ChestPain  nonanginal
RestBP      120
Chol      219
Fbs      0
```

```

RestECG          0
MaxHR             158
ExAng            0
Oldpeak          1.6
Slope            2
Ca               0
Thal             normal
AHD              No
Name: 25, dtype: object

```

```

[7]: # display of the array in the dataset.
      datanp

```

```

[7]: array([[1, 63, 1, ..., 0.0, 'fixed', 'No'],
           [2, 67, 1, ..., 3.0, 'normal', 'Yes'],
           [3, 67, 1, ..., 2.0, 'reversible', 'Yes'],
           ...,
           [301, 57, 1, ..., 1.0, 'reversible', 'Yes'],
           [302, 57, 0, ..., 1.0, 'normal', 'Yes'],
           [303, 38, 1, ..., nan, 'normal', 'No']], dtype=object)

```

2 Question 3

In this exercise, you will plot some data using the matplotlib library. Create two side-by-side plots. In one, plot the Age column against the MaxHR column in the dataset. In the other, plot the Age column against the RestBP. Set an appropriate label on the overall plot figure as well as the individual sub-plots. Finally, create a third stand-alone plot that contains both of the above plots in one, making sure that the MaxHR and RestBP plots can be clearly differentiated by changing the plot type and creating a plot legend.

Plotting 2 individual plots

```

[8]: # Plotting the 2 plots using the matplotlib library. That is Age vrs MaxHR and
      ↪Age vrs RestBP.

      plt.figure(figsize=(15,8))

      plt.subplot(1, 2, 1)
      plt.scatter('Age', 'MaxHR', data=data, color = '#88c999')
      plt.title("Age vrs MaxHR")
      plt.xlabel('Age ')
      plt.ylabel('MaxHR ')

      plt.subplot(1, 2, 2)
      plt.scatter('Age', 'RestBP', data=data, color = 'hotpink')

```

```
plt.title("Age vrs RestBP")
plt.xlabel('Age ')
plt.ylabel('RestBP ')

plt.tight_layout(3)
plt.suptitle("Plot of Age vrs MaxHR and Age vrs RestBP Plots",
    ↳ fontsize="x-large")

plt.show()
```

<ipython-input-8-d34c7d9d1397>:19: MatplotlibDeprecationWarning: Passing the pad parameter of tight_layout() positionally is deprecated since Matplotlib 3.3; the parameter will become keyword-only two minor releases later.

```
plt.tight_layout(3)
```




Plotting standalone plot with 2 plots within.

[9]: *# Codes for standalone plot showing 2 plots of scatter plot and line plot.*

```
plt.figure(figsize=(30,10))

plt.subplot(1, 2, 1)
plt.scatter('Age', 'MaxHR', data=data, color = '#88c999')
plt.plot('Age', 'RestBP', data=data)
plt.xlabel('Age')
plt.legend()
plt.title("Plot of Age vrs MaxHR and Age vrs RestBP Plots", fontsize=20)

plt.show()
```



2.1 Question 4

Use Sci-kit learn to create a simple linear regression model from the same Heart dataset. Create a predictor that can predict the MaxHR using all the other columns. First partition the data into a training and a test set. Try at least 3 different partitions based on training:test ratios of (1) 50:50, (2) 75:25 and (3) 90:10. For each partition, calculate the training error and the testing error (use simple error function based on mean error) after fitting the Linear Regression model. Provide some explanation of what is happening in terms of any changes in the training vs. testing errors across the various ratio cases.

Looking at details of data

```
[10]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   303 non-null    int64
1   Age          303 non-null    int64
```

```

2   Sex          303 non-null   int64
3   ChestPain    303 non-null   object
4   RestBP       303 non-null   int64
5   Chol         303 non-null   int64
6   Fbs          303 non-null   int64
7   RestECG      303 non-null   int64
8   MaxHR        303 non-null   int64
9   ExAng        303 non-null   int64
10  Oldpeak      303 non-null   float64
11  Slope        303 non-null   int64
12  Ca           299 non-null   float64
13  Thal         301 non-null   object
14  AHD          303 non-null   object
dtypes: float64(2), int64(10), object(3)
memory usage: 35.6+ KB

```

```
[11]: data.describe()
```

```

[11]:      Unnamed: 0      Age      Sex      RestBP      Chol      Fbs  \
count  303.000000  303.000000  303.000000  303.000000  303.000000  303.000000
mean    152.000000   54.438944   0.679868  131.689769  246.693069   0.148515
std      87.612784    9.038662   0.467299   17.599748   51.776918   0.356198
min       1.000000   29.000000   0.000000   94.000000  126.000000   0.000000
25%      76.500000   48.000000   0.000000  120.000000  211.000000   0.000000
50%     152.000000   56.000000   1.000000  130.000000  241.000000   0.000000
75%     227.500000   61.000000   1.000000  140.000000  275.000000   0.000000
max     303.000000   77.000000   1.000000  200.000000  564.000000   1.000000

      RestECG      MaxHR      ExAng      Oldpeak      Slope      Ca
count  303.000000  303.000000  303.000000  303.000000  303.000000  299.000000
mean     0.990099  149.607261   0.326733   1.039604   1.600660   0.672241
std     0.994971   22.875003   0.469794   1.161075   0.616226   0.937438
min     0.000000   71.000000   0.000000   0.000000   1.000000   0.000000
25%     0.000000  133.500000   0.000000   0.000000   1.000000   0.000000
50%     1.000000  153.000000   0.000000   0.800000   2.000000   0.000000
75%     2.000000  166.000000   1.000000   1.600000   2.000000   1.000000
max     2.000000  202.000000   1.000000   6.200000   3.000000   3.000000

```

```
[12]: data.describe().T
```

```

[12]:      count      mean      std      min      25%      50%      75%      max
Unnamed: 0  303.0  152.000000  87.612784    1.0    76.5    152.0    227.5    303.0
Age         303.0   54.438944   9.038662   29.0    48.0    56.0    61.0    77.0
Sex         303.0    0.679868   0.467299    0.0     0.0     1.0     1.0     1.0
RestBP      303.0  131.689769  17.599748   94.0   120.0   130.0   140.0   200.0
Chol        303.0  246.693069  51.776918  126.0  211.0   241.0   275.0   564.0
Fbs         303.0   0.148515   0.356198    0.0     0.0     0.0     0.0     1.0

```

| | | | | | | | | |
|---------|-------|------------|-----------|------|-------|-------|-------|-------|
| RestECG | 303.0 | 0.990099 | 0.994971 | 0.0 | 0.0 | 1.0 | 2.0 | 2.0 |
| MaxHR | 303.0 | 149.607261 | 22.875003 | 71.0 | 133.5 | 153.0 | 166.0 | 202.0 |
| ExAng | 303.0 | 0.326733 | 0.469794 | 0.0 | 0.0 | 0.0 | 1.0 | 1.0 |
| Oldpeak | 303.0 | 1.039604 | 1.161075 | 0.0 | 0.0 | 0.8 | 1.6 | 6.2 |
| Slope | 303.0 | 1.600660 | 0.616226 | 1.0 | 1.0 | 2.0 | 2.0 | 3.0 |
| Ca | 299.0 | 0.672241 | 0.937438 | 0.0 | 0.0 | 0.0 | 1.0 | 3.0 |

```
[13]: #null value check for missing values.
data.isna().sum()
```

```
[13]: Unnamed: 0      0
Age              0
Sex              0
ChestPain        0
RestBP           0
Chol             0
Fbs              0
RestECG          0
MaxHR            0
ExAng            0
Oldpeak          0
Slope            0
Ca               4
Thal             2
AHD              0
dtype: int64
```

From the above, it is clear there are missing values in the variable “Ca” and “Thal”. This number of missing data cannot be ignored. The variable “ca” needs to be corrected but the variable “Thal” can be ignored because is categorical variable which be transformed into dummy variables before building the model.

```
[14]: # Checking the variable 'Ca'.
data[data["Ca"].isnull()]
```

```
[14]: Unnamed: 0  Age  Sex  ChestPain  RestBP  Chol  Fbs  RestECG  MaxHR  \
166          167   52   1   nonanginal    138   223    0         0    169
192          193   43   1  asymptomatic    132   247    1         2    143
287          288   58   1  nontypical     125   220    0         0    144
302          303   38   1   nonanginal    138   175    0         0    173

      ExAng  Oldpeak  Slope  Ca      Thal  AHD
166      0      0.0      1 NaN    normal  No
192      1      0.1      2 NaN  reversable Yes
287      0      0.4      2 NaN  reversable No
302      0      0.0      1 NaN    normal  No
```

```
[15]: # Treating the missing values by replacing with the mean value of that column.
old_data_1 = data
updated_data = data
updated_data['Ca']=updated_data['Ca'].fillna(updated_data['Ca'].mean())
data = updated_data
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0    303 non-null    int64
1   Age          303 non-null    int64
2   Sex          303 non-null    int64
3   ChestPain    303 non-null    object
4   RestBP       303 non-null    int64
5   Chol         303 non-null    int64
6   Fbs          303 non-null    int64
7   RestECG      303 non-null    int64
8   MaxHR        303 non-null    int64
9   ExAng        303 non-null    int64
10  Oldpeak      303 non-null    float64
11  Slope        303 non-null    int64
12  Ca           303 non-null    float64
13  Thal         301 non-null    object
14  AHD          303 non-null    object
dtypes: float64(2), int64(10), object(3)
memory usage: 35.6+ KB
```

From the above, the missing values in the variable “Ca” has been corrected.

```
[16]: # Checking the uniqueness of the all variables.
data.nunique()
```

```
[16]: Unnamed: 0    303
Age            41
Sex            2
ChestPain      4
RestBP         50
Chol           152
Fbs            2
RestECG        3
MaxHR          91
ExAng          2
Oldpeak        40
Slope          3
Ca             5
```



```

Thal          3
AHD           2
dtype: int64

```

The unnamed variable is unique and can be dropped before the linear regression model is built.

```

[17]: # Dropping the variable "Unnamed: 0" as indicated above
old_data_2 = data # keeping copy of data.
data.drop(columns=["Unnamed: 0"], inplace=True)

data.nunique()

```

```

[17]: Age          41
      Sex           2
      ChestPain     4
      RestBP        50
      Chol         152
      Fbs           2
      RestECG       3
      MaxHR         91
      ExAng         2
      Oldpeak       40
      Slope         3
      Ca            5
      Thal          3
      AHD           2
      dtype: int64

```

```

[18]: data.head()

```

```

[18]:   Age  Sex  ChestPain  RestBP  Chol  Fbs  RestECG  MaxHR  ExAng  Oldpeak  \
0   63   1    typical    145    233   1         2    150     0        2.3
1   67   1  asymptomatic    160    286   0         2    108     1        1.5
2   67   1  asymptomatic    120    229   0         2    129     1        2.6
3   37   1  nonanginal    130    250   0         0    187     0        3.5
4   41   0  nontypical    130    204   0         2    172     0        1.4

      Slope  Ca      Thal  AHD
0         3  0.0    fixed   No
1         2  3.0    normal  Yes
2         2  2.0  reversable  Yes
3         3  0.0    normal   No
4         1  0.0    normal   No

```

```

[19]: data.corr()

```

```
[19]:
```

| | Age | Sex | RestBP | Chol | Fbs | RestECG | MaxHR | \ |
|---------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|---|
| Age | 1.000000 | -0.097542 | 0.284946 | 0.208950 | 0.118530 | 0.148868 | -0.393806 | |
| Sex | -0.097542 | 1.000000 | -0.064456 | -0.199915 | 0.047862 | 0.021647 | -0.048663 | |
| RestBP | 0.284946 | -0.064456 | 1.000000 | 0.130120 | 0.175340 | 0.146560 | -0.045351 | |
| Chol | 0.208950 | -0.199915 | 0.130120 | 1.000000 | 0.009841 | 0.171043 | -0.003432 | |
| Fbs | 0.118530 | 0.047862 | 0.175340 | 0.009841 | 1.000000 | 0.069564 | -0.007854 | |
| RestECG | 0.148868 | 0.021647 | 0.146560 | 0.171043 | 0.069564 | 1.000000 | -0.083389 | |
| MaxHR | -0.393806 | -0.048663 | -0.045351 | -0.003432 | -0.007854 | -0.083389 | 1.000000 | |
| ExAng | 0.091661 | 0.146201 | 0.064762 | 0.061310 | 0.025665 | 0.084867 | -0.378103 | |
| Oldpeak | 0.203805 | 0.102173 | 0.189171 | 0.046564 | 0.005747 | 0.114133 | -0.343085 | |
| Slope | 0.161770 | 0.037533 | 0.117382 | -0.004062 | 0.059894 | 0.133946 | -0.385601 | |
| Ca | 0.359489 | 0.092891 | 0.098707 | 0.118525 | 0.143967 | 0.127487 | -0.263408 | |

| | ExAng | Oldpeak | Slope | Ca |
|---------|-----------|-----------|-----------|-----------|
| Age | 0.091661 | 0.203805 | 0.161770 | 0.359489 |
| Sex | 0.146201 | 0.102173 | 0.037533 | 0.092891 |
| RestBP | 0.064762 | 0.189171 | 0.117382 | 0.098707 |
| Chol | 0.061310 | 0.046564 | -0.004062 | 0.118525 |
| Fbs | 0.025665 | 0.005747 | 0.059894 | 0.143967 |
| RestECG | 0.084867 | 0.114133 | 0.133946 | 0.127487 |
| MaxHR | -0.378103 | -0.343085 | -0.385601 | -0.263408 |
| ExAng | 1.000000 | 0.288223 | 0.257748 | 0.144722 |
| Oldpeak | 0.288223 | 1.000000 | 0.577537 | 0.294558 |
| Slope | 0.257748 | 0.577537 | 1.000000 | 0.109618 |
| Ca | 0.144722 | 0.294558 | 0.109618 | 1.000000 |

2.2 Linear Regression Model Building

Divide the data into independent and dependent variables

```
[20]: # Divide the data into independent and dependent variables
ind_vars = data.drop(["MaxHR"], axis=1) # independent variables
dep_var = data[["MaxHR"]] # dependent variables

ind_vars.head() # Print
```

```
[20]:
```

| | Age | Sex | ChestPain | RestBP | Chol | Fbs | RestECG | ExAng | Oldpeak | Slope | \ |
|---|-----|-----|--------------|--------|------|-----|---------|-------|---------|-------|---|
| 0 | 63 | 1 | typical | 145 | 233 | 1 | 2 | 0 | 2.3 | 3 | |
| 1 | 67 | 1 | asymptomatic | 160 | 286 | 0 | 2 | 1 | 1.5 | 2 | |
| 2 | 67 | 1 | asymptomatic | 120 | 229 | 0 | 2 | 1 | 2.6 | 2 | |
| 3 | 37 | 1 | nonanginal | 130 | 250 | 0 | 0 | 0 | 3.5 | 3 | |
| 4 | 41 | 0 | nontypical | 130 | 204 | 0 | 2 | 0 | 1.4 | 1 | |

| | Ca | Thal | AHD |
|---|-----|------------|-----|
| 0 | 0.0 | fixed | No |
| 1 | 3.0 | normal | Yes |
| 2 | 2.0 | reversible | Yes |
| 3 | 0.0 | normal | No |

```
4  0.0      normal    No
```

Above is display of independent variables.

```
[21]: dep_var.head()  # Print
```

```
[21]:    MaxHR
0     150
1     108
2     129
3     187
4     172
```

Above is the display of independent variables.

There are three (3) categorical variables (namely - ChestPain, Thal and AHD) that needs to be converted to dummy variables before the modelling can be then for data.

```
[22]: # Creating dummy variables for the 3 categorical variables.
def encode_cat_vars(x):
    x = pd.get_dummies(
        x,
        columns=x.select_dtypes(include=["object", "category"]).columns.
        →tolist(),
        drop_first=True,
    )
    return x

ind_vars_num = encode_cat_vars(ind_vars)
ind_vars_num.head()
```

```
[22]:    Age  Sex  RestBP  Chol  Fbs  RestECG  ExAng  Oldpeak  Slope  Ca  \
0    63    1    145   233    1         2     0     2.3     3  0.0
1    67    1    160   286    0         2     1     1.5     2  3.0
2    67    1    120   229    0         2     1     2.6     2  2.0
3    37    1    130   250    0         0     0     3.5     3  0.0
4    41    0    130   204    0         2     0     1.4     1  0.0

    ChestPain_nonanginal  ChestPain_nontypical  ChestPain_typical  Thal_normal  \
0                      0                      0                  1              0
1                      0                      0                  0              1
2                      0                      0                  0              0
3                      1                      0                  0              1
4                      0                      1                  0              1

    Thal_reversable  AHD_Yes
0                  0        0
```

| | | |
|---|---|---|
| 1 | 0 | 1 |
| 2 | 1 | 1 |
| 3 | 0 | 0 |
| 4 | 0 | 0 |

Above is the display of 6 dummy variables created for the data.

```
[23]: dep_var.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0   MaxHR    303 non-null      int64
dtypes: int64(1)
memory usage: 2.5 KB
```

2.3 Part 1

2.3.1 First splitting of data into training (50%) and test data (50%) sets

```
[24]: # Splitting data into training (50%) and test data (50%) sets.
```

```
x_train, x_test, y_train, y_test = train_test_split(
    ind_vars_num, dep_var, test_size=0.5, random_state=1)
```

```
[25]: # Shape of the train and test sets
```

```
print("Number of rows in train data =", x_train.shape[0])
print("Number of rows in test data =", x_test.shape[0])
print("Number of rows in train data =", y_train.shape[0])
print("Number of rows in test data =", y_test.shape[0])
```

```
Number of rows in train data = 151
```

```
Number of rows in test data = 152
```

```
Number of rows in train data = 151
```

```
Number of rows in test data = 152
```

Fitting a linear regression model

```
[26]: # Fitting the linear regression model for the training (50%) and test data,
      ↪ (50%) sets
```

```
regressor = LinearRegression()
regressor.fit(x_train, y_train)
```

```
[26]: LinearRegression()
```

```
[27]: print ("The intercept is", regressor.intercept_)
```

The intercept is [145.08629771]

```
[28]: print ('The coefficient are', regressor.coef_)
```

The coefficient are [[-0.52702166 3.86019681 0.13221783 0.03710454
1.26097622 -0.44519288
-4.45441651 -0.25300197 -5.2907283 -2.13259871 2.77478326 6.15211424
1.55829588 18.87048387 14.09293057 -4.39477778]]

```
[29]: y_pred = regressor.predict(x_test)  
print ('The predicted value is ', y_pred)
```

The predicted value is [[157.45978472]
[148.18006805]
[158.84297735]
[121.61703462]
[153.9668921]
[168.76202868]
[167.03325878]
[137.34935307]
[139.84247603]
[130.78048063]
[145.82473123]
[161.01523843]
[160.19016662]
[138.519901]
[166.60188049]
[168.57994551]
[149.76768993]
[159.24056003]
[163.48526299]
[139.09166996]
[158.21416902]
[133.1813511]
[157.45691203]
[144.28190259]
[149.83317244]
[162.85644119]
[138.83791439]
[148.88413761]
[142.41995843]
[134.47722498]
[145.28967041]
[127.55068471]
[141.87193763]
[153.80533339]]

[141.19305756]
[123.44308342]
[158.72277511]
[151.26667971]
[134.614343]
[145.74731749]
[153.76939217]
[148.01443175]
[147.97105156]
[137.63169834]
[126.94249644]
[165.21339822]
[153.35492403]
[166.72333509]
[173.48725326]
[134.13097012]
[163.4285405]
[163.91028745]
[134.56183364]
[160.23292537]
[146.41543603]
[156.77204251]
[129.03587005]
[158.88563359]
[155.46919396]
[157.65771453]
[137.01954764]
[130.13519462]
[135.72054675]
[142.57376321]
[151.19823338]
[133.93036719]
[168.45175776]
[158.28143432]
[146.39493611]
[148.31091929]
[153.26461929]
[147.90956682]
[159.46193357]
[143.09437604]
[117.02815435]
[161.14195941]
[137.76171093]
[146.47608828]
[156.68628568]
[155.39365541]
[148.53221644]
[177.26199462]

[143.7273677]
[163.58556088]
[152.32061981]
[142.73703817]
[140.75062562]
[161.18647001]
[159.21004934]
[165.67088025]
[145.08803296]
[163.39013882]
[153.09648702]
[124.76930186]
[167.77767736]
[154.06886904]
[132.20890353]
[136.59508552]
[169.76062745]
[164.37809923]
[124.62012593]
[156.27771905]
[145.6767857]
[150.6653847]
[156.55658791]
[140.36483209]
[151.6323045]
[157.11311193]
[159.31180674]
[147.09035927]
[121.31093072]
[145.82398458]
[144.08838994]
[164.90228387]
[125.49285768]
[160.70660786]
[142.39800805]
[138.09651001]
[144.13782419]
[153.63691769]
[142.73794417]
[137.92330928]
[141.39609272]
[137.83483441]
[168.48231175]
[160.01143121]
[152.63567908]
[159.20310839]
[141.9372064]
[147.76905901]

```
[167.28746271]
[162.00532295]
[163.37183045]
[160.69927135]
[160.33252397]
[147.71967209]
[160.61305617]
[158.51342725]
[134.97859203]
[156.34043555]
[149.01299464]
[170.11598715]
[148.61612796]
[130.58659818]
[160.10794079]
[167.63015407]
[158.86554715]
[169.91323778]
[144.45529022]
[131.05653985]
[161.47589756]
[157.47141883]
```

```
[42]: # actual_pred = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
      # print(actual_pred)
```

Evaluate the algorithm

```
[30]: # Evaluating or calculating the metrics and displaying them.
      print ('Mean Absolute error:', metrics.mean_absolute_error(y_test, y_pred))
      print ('Mean squared error:', metrics.mean_squared_error(y_test, y_pred, ↵
        ↪squared=True))
      print ('Root Mean squared error:', metrics.mean_squared_error(y_test, y_pred, ↵
        ↪squared=False))
```

Mean Absolute error: 15.0254336413446

Mean squared error: 376.76275647699526

Root Mean squared error: 19.410377545967396

```
[44]: # Keeping results in a variable
      mae5050 = metrics.mean_absolute_error(y_test, y_pred)
      mse5050 = metrics.mean_squared_error(y_test, y_pred, squared=True)
      rmse5050 = metrics.mean_squared_error(y_test, y_pred, squared=False)
```


2.4 Part 2

2.4.1 First splitting of data into training (75%) and test data (25%) sets

```
[47]: # Splitting data into training (75%) and test data (25%) sets.
```

```
x_train, x_test, y_train, y_test = train_test_split(  
    ind_vars_num, dep_var, test_size=0.25, random_state=1)
```

```
[48]: # Shape of the train and test sets
```

```
print("Number of rows in train data =", x_train.shape[0])  
print("Number of rows in test data =", x_test.shape[0])  
print("Number of rows in train data =", y_train.shape[0])  
print("Number of rows in test data =", y_test.shape[0])
```

Number of rows in train data = 227

Number of rows in test data = 76

Number of rows in train data = 227

Number of rows in test data = 76

Fitting a linear regression model

```
[49]: # Fitting the linear regression model for the training (75%) and test data  
      ↪ (25%) sets
```

```
regressor = LinearRegression()  
regressor.fit(x_train, y_train)
```

```
[49]: LinearRegression()
```

```
[50]: print ("The intercept is", regressor.intercept_)
```

The intercept is [172.66013717]

```
[51]: print ('The coefficient are', regressor.coef_)
```

```
The coefficient are [[-0.87167798  1.43477136  0.1510971  0.01825345  
5.51742771 -0.95452007  
-5.52181268  0.9612542  -7.87616602  0.18432209  2.7989238  4.13287437  
6.23820165 16.77397426 13.33148954 -8.00491205]]
```

```
[52]: y_pred = regressor.predict(x_test)  
print ('The predicted value is ', y_pred)
```

The predicted value is [[162.54023894]
[147.1062943]
[155.45276195]
[120.71119829]
[148.78396003]

[172.38423856]
[172.10941047]
[133.01136511]
[145.74752509]
[120.93831243]
[144.63817874]
[159.82709153]
[155.21407129]
[131.09926047]
[163.60833047]
[176.83754076]
[145.59493735]
[163.00202027]
[167.37090011]
[133.17538984]
[154.86138695]
[131.83334945]
[156.79703387]
[149.31238722]
[148.25356693]
[173.22784084]
[144.87512176]
[150.85476848]
[143.80217526]
[132.82650375]
[143.81077512]
[127.48577837]
[142.96269844]
[161.64969009]
[135.08556112]
[121.73865615]
[167.66223427]
[152.37688356]
[132.67870333]
[144.87563429]
[155.0188081]
[142.34981725]
[150.92245487]
[137.97197848]
[131.83948235]
[178.66526455]
[145.07177346]
[172.75506625]
[178.55537149]
[138.56908359]
[159.80681229]
[161.05923653]
[138.83028693]

```
[162.78838669]
[139.84729568]
[158.5726807 ]
[133.77030626]
[156.63673605]
[149.78111792]
[167.88703444]
[130.10531738]
[131.63187899]
[134.93582095]
[141.15432166]
[148.66092576]
[131.49737201]
[170.97192532]
[162.09291805]
[152.47550352]
[144.13768493]
[152.8001417 ]
[150.90274992]
[165.31506106]
[138.96268662]
[118.76180818]
[162.59178073]]
```

Evaluate the algorithm

```
[55]: # Evaluating or calculating the metrics and displaying them.
print ('Mean Absolute error:', metrics.mean_absolute_error(y_test, y_pred))
print ('Mean squared error:', metrics.mean_squared_error(y_test, y_pred, ↵
↪squared=True))
print ('Root Mean squared error:', metrics.mean_squared_error(y_test, y_pred, ↵
↪squared=False))
```

```
Mean Absolute error: 14.557058704590277
Mean squared error: 359.87057126167997
Root Mean squared error: 18.970254907662152
```

```
[56]: # Keeping results in a variable
mae7525 = metrics.mean_absolute_error(y_test, y_pred)
mse7525 = metrics.mean_squared_error(y_test, y_pred, squared=True)
rmse7525 = metrics.mean_squared_error(y_test, y_pred, squared=False)
```

2.5 Part 3

2.5.1 First splitting of data into training (90%) and test data (10%) sets

```
[57]: # Splitting data into training (90%) and test data (10%) sets.
```

```
x_train, x_test, y_train, y_test = train_test_split(  
    ind_vars_num, dep_var, test_size=0.10, random_state=1)
```

```
[58]: # Shape of the train and test sets
```

```
print("Number of rows in train data =", x_train.shape[0])  
print("Number of rows in test data =", x_test.shape[0])  
print("Number of rows in train data =", y_train.shape[0])  
print("Number of rows in test data =", y_test.shape[0])
```

Number of rows in train data = 272

Number of rows in test data = 31

Number of rows in train data = 272

Number of rows in test data = 31

Fitting a linear regression model

```
[59]: # Fitting the linear regression model for the training (75%) and test data  
      ↪ (25%) sets
```

```
regressor = LinearRegression()  
regressor.fit(x_train, y_train)
```

```
[59]: LinearRegression()
```

```
[60]: print ("The intercept is", regressor.intercept_)
```

The intercept is [170.84925848]

```
[61]: print ('The coefficient are', regressor.coef_)
```

The coefficient are [[-0.82496715 1.74109573 0.13827588 0.03215715
3.82255152 -0.25874941
-8.80188544 -0.37930585 -5.74395412 0.08906314 3.29001685 7.00779945
8.28576312 10.79653614 8.72227747 -6.30072826]]

```
[62]: y_pred = regressor.predict(x_test)  
print ('The predicted value is ', y_pred)
```

The predicted value is [[162.0905965]
[147.69477114]
[156.24903587]
[127.77315533]
[151.24110557]

```
[172.46886147]
[170.04848505]
[135.24579945]
[144.61452951]
[121.49537377]
[144.07901254]
[161.19582594]
[159.65338962]
[129.93033974]
[165.67858131]
[174.3536468 ]
[150.34211827]
[160.15383582]
[166.33665483]
[133.65442728]
[158.76975559]
[129.52133012]
[155.34858394]
[151.80274175]
[144.82509814]
[170.68149118]
[140.26434514]
[145.3875382 ]
[145.27456729]
[127.81757541]
[148.01211334]]
```

Evaluate the algorithm

```
[63]: # Evaluating or calculating the metrics and displaying them.
print ('Mean Absolute error:', metrics.mean_absolute_error(y_test, y_pred))
print ('Mean squared error:', metrics.mean_squared_error(y_test, y_pred,
↪squared=True))
print ('Root Mean squared error:', metrics.mean_squared_error(y_test, y_pred,
↪squared=False))
```

```
Mean Absolute error: 13.220295962205572
Mean squared error: 286.01884104920157
Root Mean squared error: 16.91209156341112
```

```
[64]: # Keeping results in a variable
mae9010 = metrics.mean_absolute_error(y_test, y_pred)
mse9010 = metrics.mean_squared_error(y_test, y_pred, squared=True)
rmse9010 = metrics.mean_squared_error(y_test, y_pred, squared=False)
```

```
[89]: # Printing metrics to compare results
print ("Ratio", "MAE", "      MSE      ", " RMSE      ")
print ("")
```

```
print ("50:50", round(mae5050, 5), round(mse5050, 5), round(rmse5050, 5))
print ("75:25", round(mae7525, 5), round(mse7525, 5), round(rmse7525, 5))
print ("90:10", round(mae9010, 5), round(mse9010, 5), round(rmse9010, 5))
```

| Ratio | MAE | MSE | RMSE |
|-------|----------|-----------|----------|
| 50:50 | 15.02543 | 376.76276 | 19.41038 |
| 75:25 | 14.55706 | 359.87057 | 18.97025 |
| 90:10 | 13.2203 | 286.01884 | 16.91209 |

2.5.2 Explanation on changes in errors after changes in the ratios

The ratios provided and used in the 3 linear regression models varied the amount of data that was used to build and test the models. In the first model in part 1 above, it was 50% split, the other 2 parts had more training data than the test data as in 75:25 and 90:10. From the list of errors obtained, the errors were reducing as the training data was increased. This shows that more training data leads to reduced errors. The lesser the mean square error, the better the regression model is. When the linear regression model is trained using a given set of observations, the model with the least mean sum of mean square error is selected as the best model. In this exercise, the model with ratio of 90:10 is the model with least error close to zero. This will taken as the best model.